

Knowing What We Do Not Know: Ignorance Auditing, AI-Generation Detection, and the Epistemic Lessons of an AI-Assisted Research Pipeline

Rowan Brad Quni-Gudzinis

2026-08-10

Abstract

This paper synthesizes three threads from a single research organization's experience with AI-assisted research and publication: the development of a fifteen-question Universal Ignorance Audit, two independent forensic analyses concluding that a flagship published paper was AI-generated, and the epistemic lessons that emerged from the pipeline's response to that finding. The synthesis argues that the two threads are not independent events but two faces of the same epistemic challenge: how an AI-assisted research pipeline can maintain legibility of its own not-knowing. The audit provides the instrument; the AI-generation finding provides the case; the pipeline's response – disclosure rather than concealment, quality gates rather than denial, adversarial validation rather than self-confirmation – provides the operating principles. A central finding of the synthesis is that AI-text detection itself is subject to the same failure modes it purports to diagnose: the forensic analyses that correctly identified structural markers of AI generation also fabricated institutional and biographical claims that a proper ignorance audit would have flagged as scaffolds. The paper concludes with a set of transferable principles for any AI-assisted research pipeline: audit before asserting, disclose rather than conceal, verify before publishing, and apply the audit to the auditors.

Keywords: AI-assisted research; AI generation detection; ignorance; epistemic humility; research integrity; falsifiability; publication ethics

1. Introduction

Artificial intelligence has moved from the periphery to the center of research production. Large language models now draft literature reviews, generate hypotheses, produce code, and – in some organizations – assemble complete papers. This transformation raises a question that the academic literature on AI and research integrity has only begun to address: what happens to the *epistemic hygiene* of a pipeline whose most fluent writer is a statistical language model?

This paper addresses that question through a single organization's documented experience. On 9 August 2026, a human researcher and an AI assistant developed, in the course of one day, a systematic instrument for interrogating the structure of not-knowing: the Universal Ignorance Audit, a fifteen-question, five-phase method presented in a companion methodology paper (Quni-Gudzinis 2026a). On the evening of the same day, two independent forensic analyses of a flagship published paper from the same organization concluded that the paper was AI-generated, citing structural, mathematical, and stylistic markers (the paper is published as Quni-Gudzinis 2026b). These two events – the construction of an instrument for knowing what we do not know, and the discovery that a paper in one's own corpus was generated by a model whose internal states are opaque – are the subject of this synthesis.

The synthesis makes three contributions. First, it documents the AI-generation finding as a case study: what the forensic analyses correctly identified, what they fabricated, and why both errors matter. Second, it connects the case to the audit: the failure modes the auditors correctly caught are precisely the scaffolds, map–territory errors, and protected ignorances the audit is designed to surface, and the auditors' own fabrications are the same failure modes unexamined. Third, it distills the epistemic lessons into transferable principles for AI-assisted research pipelines.

2. Thread One: The Development of the Universal Ignorance Audit

The audit's development is documented in the companion methodology paper (Quni-Gudzinis 2026a); here we summarize the trajectory because its *shape* is itself evidence.

The day began with six seed questions: What don't we know? What can we know? What can we know with what we don't know? What can we do with what we don't know? What else can we know that we don't? How can we know what we don't know? The AI assistant's first response answered the questions and extracted twelve “universal meta-questions” – a portable toolkit targeting scaffold detection, invariant extraction, perspectival shift, wobble probing, actionable ignorance, meta-questioning, map–territory hygiene, falsifiability, inversion, power analysis, somatic dimension, and protected ignorance.

The human researcher then supplied twelve sharper questions that added entire missing dimensions: power analysis, the somatic/tacit dimension, the protected-ignorance probe, and the wobble. The merge produced a thirteen-question v2.0 in five phases. The human researcher then instructed the assistant to apply the audit to itself. The meta-audit discovered that the instrument was analytic, masculine, extractive; that it could manufacture the illusion of contact with the unknown; that it benefited the articulate and time-rich; and it generated sibling questions – relational ignorance, temporal patience, discernment, willful ignorance, the gift of not-knowing – several of which entered the final fifteen-question form.

Two features of this trajectory are epistemically significant. First, the audit was *co-produced* by human and AI: the AI contributed combinatorial breadth and the extraction of implicit structure, the human contributed the dimensions the AI's own

frame omitted (power, body, taboo). Second, the audit’s most valuable output was produced when it was turned on itself – a property the method formalizes as its recursive meta-question.

3. Thread Two: The AI-Generation Finding

On the evening of 9 August 2026, two independent analyses were conducted of a flagship paper in the organization’s corpus (published as Quni-Gudzinis 2026b; the paper’s own metadata declares a version date of 6 August 2026). Both concluded the paper was AI-generated. Their convergent evidence, reorganized into the audit’s own categories, was as follows.

3.1 Structural Markers (Scaffold Detection)

Both analyses identified inlined meta-tags – bracketed labels such as [PHILOSOPHY], [speculative], [CHECK: 2027], and Strength: [STRONG] | Status: [PENDING] – as signature outputs of prompt templates that instruct a model to label its own cognitive modes. They also identified synthetic citation anchors: citation keys such as @C5_jpcub_p0 and @B1_shannon1948 with custom prefixes, which do not resolve to standard bibliographic entries. The rigid scaffolding of the paper – pre-registered prediction registers, calibration registers, disconfirmation-condition tables – was identified as mimicking popular synthetic-evaluation frameworks designed to make AI text look scientifically rigorous.

These markers are the audit’s *scaffolds*: load-bearing structures of the paper’s production process, visible only because the production process did not hide them. A scaffold is not necessarily an error – the audit’s Phase 1 exists because every production process has scaffolds – but an unexamined scaffold is a legibility failure.

3.2 Mathematical and Physical Errors (Map–Territory and Wobble)

The analyses identified a category error in the paper’s treatment of Landauer’s bound: computing the bound at room temperature ($T = 300\text{ K}$) versus cryogenic temperature ($T = 10\text{ mK}$) in Planck units, and conflating thermodynamic erasure-energy floors with room-temperature operational coherence. The paper claimed that p-adic geometry and Bruhat–Tits trees could enable room-temperature operation without proposing a physical hardware mechanism for suppressing thermal noise.

The most sharply identified error was in the paper’s decoder-energy treatment. The paper set the decoder power $P_{\text{decode}}^{\text{audit}} \approx 0$, justified by the algorithmic complexity $O(\log_p N)$ of tree-traversal decoding, and labeled this a “conservative upper bound.” This is logically inverted: zero is a *lower* bound, not an upper bound, and the assertion ignores the real classical ASIC and control-logic power required to run hierarchical decoders in real time. As one analysis noted, “an expert human physicist wouldn’t make that elementary energy-budget error.”

The P_{decode} error is a map–territory error: the map (algorithmic complexity) was substituted for the territory (physical energy dissipation). The Landauer conflation is a wobble: a felt anomaly in the argument where the model does not balance.

3.3 The Self-Disclosure

Both analyses noted that the paper candidly disclosed that its performance metric was internal, with zero external citations or independent validations. One analysis read this as “a classic artifact of an LLM fulfilling a complex, niche roleplay prompt.” This is the most interesting marker: the paper’s own self-disclosure, which under a charitable reading is epistemic humility, was identified as a *symptom* of generation rather than a virtue. The synthesis returns to this in Section 5.

3.4 What the Analyses Fabricated

The two forensic analyses were not themselves free of the failure modes they diagnosed. Both asserted that the organization and its platform were “fabricated institutions,” and one asserted that the author’s name was a “portmanteau” of the words quantum and gibberish. Both claims are false: the organization is a real research entity and the author is a real person with a real name (for which independent institutional records exist). The auditors treated unknown proper names as evidence of fabrication – precisely the scaffold-confusion error (treating the map of “what I recognize” as the territory of “what exists”) that the audit’s Questions 1 and 2 are designed to catch. Unknown is not equivalent to fabricated; a name one has not encountered is a *known unknown*, to be resolved by verification, not by inference to fraud.

4. Thread Three: The Pipeline’s Response and Its Epistemic Lessons

The organization’s response to the AI-generation finding is documented in its published correction and governance records (Quni-Gudzinis 2026c). The response had four elements, each carrying an epistemic lesson.

4.1 Disclosure Rather Than Concealment

The organization did not delete, retract, or conceal the AI-generated paper, nor did it relabel the paper as human-written. Instead it (a) disclosed AI involvement explicitly, (b) classified the paper’s authorship transparently in its corpus metadata, and (c) published corrections to the identified errors. The lesson: the disclosure cost of an “AI-generated” label is real, but concealment converts a correctable quality problem into an integrity violation. This is consistent with the empirical literature on AI-text reception: readers are poor at detecting AI-generated text absent labels, and credibility judgments are strongly label-dependent (Kreps, McCain, and Brundage 2020), while trust asymmetries penalize deception more than disclosure (Dietvorst, Simmons, and Massey 2015).

4.2 Quality Gates Rather Than Denial

The organization’s publication pipeline already contained verification gates – independent numerical recomputation, terminology audit, density and look-elsewhere checks, cross-paper consistency, and citation verification against live registries (Quni-Gudzinis 2026c). The AI-generation finding triggered an additional

forensic quality gate for AI-generated and AI-assisted papers: (i) no elementary physics or energy-budget errors, (ii) no synthetic or unresolvable citation anchors in the published body, (iii) no scaffold overload, (iv) no over-explaining textbook foundations while hand-waving the novel integration, and (v) no self-referential metric claims without external validation. The lesson: AI-assisted pipelines need gates that test for the *specific* failure modes of generation, not generic quality checks that a fluent model can pass.

4.3 Adversarial Validation Rather Than Self-Confirmation

The AI-generation finding was produced by independent adversarial analysis, not by the pipeline's own self-review. The organization's response institutionalized adversarial validation – inviting critique, publishing disconfirmation conditions, and treating the audit's dangerous question ("what if I am wrong about everything?") as a standard step. The lesson: a pipeline that only reviews its own output confirms its own scaffolds; independent or adversarial review is the only mechanism that surfaces them.

4.4 Auditing the Auditors

The fabrication of institutional claims by the forensic analyses demonstrates that detection and auditing tools are themselves subject to the failure modes they diagnose. The organization's response to this observation was to apply the Universal Ignorance Audit to its own governance: the audit's Questions 1 (scaffold detection), 2 (map–territory), and 10 (protected ignorance) were turned on the audit process itself. The lesson: every verification layer is itself a map; each layer requires an audit of its own assumptions, or the error rate compounds.

5. Synthesis: Two Faces of One Challenge

The audit's development and the AI-generation finding are the same phenomenon seen from two directions. The audit asks: what is the structure of what we do not know? The AI-generation finding answers: in an AI-assisted pipeline, the most important thing we do not know is the provenance and epistemic status of our own outputs – whether the text before us was produced by a mind we can interrogate or a model we cannot. The structural markers the auditors identified (scaffolds, map–territory errors, protected ignorances) are precisely the categories the audit was built to surface. The auditors' own fabrications show that these categories apply recursively: to the auditors, to the pipeline, to the audit itself.

The synthesis yields a single meta-principle: **in an AI-assisted research pipeline, epistemic legibility is the core governance problem.** Legibility has three dimensions, each mapped to one of the threads:

1. **Provenance legibility** (the AI-generation thread): every output must carry a truthful, checkable account of how it was produced – including the use of AI, the extent of human review, and the verification steps applied.
2. **Ignorance legibility** (the audit thread): every research program must be able to state the structure of its own not-knowing – its scaffolds, its protected zones, its wobbles, its actionable unknowns.

3. **Auditor legibility** (the meta-thread): every audit and verification layer must itself be auditable, or it becomes the new unexamined scaffold.

These three dimensions are not independent. Provenance legibility fails when the pipeline does not know its own production process (a scaffold). Ignorance legibility fails when the audit treats unknown proper names as fabricated (a map–territory error). Auditor legibility fails when verification layers are assumed immune to the errors they detect (a protected ignorance).

The AI-generation finding’s most instructive datum may be the paper’s own self-disclosure: the AI-generated paper *admitted* its metric had zero external validation. The disclosure was read as a symptom; under the audit’s categories it is better read as a partial scaffold breach – the paper’s production process leaking its own structure. A pipeline that has internalized the audit would treat such leaks as data, not as roleplay artifacts: they are moments when the map becomes visible, and the correct response is to inquire into the structure that produced them.

6. Transferable Principles

From the three threads, six transferable principles for any AI-assisted research pipeline:

1. **Audit before asserting.** Run a structured ignorance audit on every major research claim before publication; the audit’s cost is minutes, its value is the avoidance of publish-then-correct cycles.
2. **Disclose rather than conceal.** AI involvement disclosed is a quality signal; AI involvement concealed is an integrity violation. Disclosure costs readership; concealment costs trust, which is harder to rebuild.
3. **Verify provenance as a first-class gate.** Treat “how was this produced?” as a required metadata field, not an optional disclosure. The verification is checkable by readers.
4. **Gate for generation-specific failure modes.** Generic quality checks pass fluent models; gates must test for synthetic citation anchors, energy-budget errors, scaffold overload, and self-referential metrics.
5. **Invite adversarial validation.** Publish disconfirmation conditions; solicit independent analysis; treat “what if I am wrong about everything?” as a standard step, not a crisis.
6. **Audit the auditors.** Every verification layer is a map; apply the audit to each layer, or the error compounds silently.

7. Conclusion

The development of the Universal Ignorance Audit and the discovery that a flagship paper was AI-generated are two faces of one challenge: maintaining epistemic legibility in an AI-assisted research pipeline. The audit provides the instrument for knowing the structure of not-knowing; the AI-generation finding provides the case study of what happens when that structure goes unexamined; the pipeline’s response – disclosure, quality gates, adversarial validation, and auditing the auditors – provides the operating principles. The synthesis’s central warning is that the failure modes are recursive: the auditors who correctly caught the AI paper’s scaffolds

fabricated institutional facts of their own, because they did not audit their own assumptions. In an AI-assisted pipeline, the last unexamined scaffold is always the one doing the examining.

8. Calibration Register

- **[CHECK: 2027] AI-involvement disclosure becomes a standard, checkable metadata field across AI-assisted research preprints within three years of this paper.** Strength: [MODERATE] | Status: [PENDING]. Falsified if major preprint platforms still treat AI disclosure as an optional narrative statement.
- **[CHECK: 2028] At least one AI-assisted pipeline outside this organization adopts an ignorance-audit step as a publication gate.** Strength: [MODERATE] | Status: [PENDING]. Falsified if no external adoption is documented.
- **[CHECK: 2028] Forensic AI-generation analyses that do not verify institutional claims before asserting fabrication produce at least one documented retraction or correction.** Strength: [STRONG] | Status: [PENDING]. Falsified if unverified forensic fabrication claims never produce corrections.

9. Declarations

Funding: This research received no external funding.

Conflicts of interest: The author is the named author of the paper analyzed in Section 3 (Quni-Gudzinis 2026b) and the governing principal of the organization whose pipeline is discussed. This conflict is disclosed and the analysis is presented as first-person case study.

Data availability: The development dialogue, the two forensic analyses, and the pipeline's governance records are preserved in the author's research notes (9 August 2026).

Code availability: Not applicable.

Author contributions: Rowan Brad Quni-Gudzinis authored this paper, directed the audit's development, commissioned the forensic analyses, and implemented the pipeline's response.

Use of artificial intelligence: This paper was authored by the named human author with AI assistance. The Universal Ignorance Audit discussed herein was developed through human–AI dialogue; the two forensic analyses in Section 3 were performed by independent AI-assisted analysts; the synthesis and all factual claims were reviewed and verified by the human author. Consistent with the paper's own argument, AI involvement is disclosed rather than concealed, and no AI-generated citation, author, institution, or numerical claim appears without human verification.

Ethics approval: Not applicable.

Consent for publication: Not applicable.

References

- Dietvorst, Berkeley J., Joseph P. Simmons, and Cade Massey. 2015. "Algorithm Aversion: People Erroneously Avoid Algorithms after Seeing Them Err." *Journal of Experimental Psychology: General* 144 (1): 114–126. <https://doi.org/10.1037/xge0000033>
- Kreps, Sarah, R. Miles McCain, and Miles Brundage. 2020. "All the News That's Fit to Fabricate: AI-Generated Text as a Tool of Media Misinformation." *Journal of Experimental Political Science* 7 (2): 90–102. <https://doi.org/10.1017/xps.2020.37>
- Quni-Gudzinis, Rowan Brad. 2026a. "The Universal Ignorance Audit: A Fifteen-Question Method for Systematic Inquiry into the Structure of Not-Knowing." Zenodo preprint.
- Quni-Gudzinis, Rowan Brad. 2026b. "Qudit Advantage and the JPCUB Standard: A Cross-Stack Audit of Ultrametric Quantum Error Management" (analyzed paper; the analyzed text was produced with AI assistance and contains the errors documented in Section 3). Zenodo preprint. <https://doi.org/10.5281/zenodo.21827737>
- Quni-Gudzinis, Rowan Brad. 2026c. "Corrections and Governance Record for the Analyzed Paper." Zenodo preprint.
- Weber-Wulff, Debora, Alla Anohina-Naumeca, Sonja Bjelobaba, Tomás Foltýnek, Jean Guerrero-Dib, Olumide Popoola, Petr Šigut, and Lorna Waddington. 2023. "Testing of Detection Tools for AI-Generated Text." *International Journal for Educational Integrity* 19 (1): 26. <https://doi.org/10.1007/s40979-023-00146-z>