

The Purpose Principle: Grounding Research in Reality, Utility, and Premise-Depth

Rowan Brad Quni-Gudzinis

2026-08-16

Abstract

Why a reader should care (SO-WHAT-GATE-1): Most research effort today is allocated by disciplinary prestige and publication counts, neither of which measures whether the work has any traceable connection to a human problem. As AI pipelines make it possible to generate silo-fodder at arbitrary scale, the question “should this research be advanced?” has become urgent and unanswerable with existing tools. This paper proposes a *grounding criterion* — a usable, falsifiable test for whether a research program deserves advancement — and a premise-depth disclosure standard that bounds novelty claims by the depth of their assumptions.

Premise-depth disclosure: The argument is as deep as three premises and no deeper: L0, a value primitive (real-world utility is a legitimate first-order criterion for allocating research effort — asserted, not derived); L1, an empirical premise (disciplinary silos, abstraction-for-its-own-sake, and knowledge siloes causally produce low-traceable-utility research — falsifiable via traceable-impact audits); L2, an epistemic premise (declaring where premises end increases testability and accountability — falsifiable, with a self-application meta-test). Everything else in this paper is derived from these three.

Falsification conditions: F1 — if a traceable-impact audit finds siloed, abstraction-only programs producing real-world utility at rates statistically indistinguishable from grounded programs, L1 fails. F2 — if requiring premise-depth disclosure measurably reduces research output quality or downstream utility, L2 fails. F3 (self-siloing meta-test) — if applying the criterion to itself reduces total utility, the criterion is self-refuting and must be revised or withdrawn.

Criterion (G1–G3): A research program’s claim to advancement should be evaluated by whether it (G1) exhibits a live, non-vacuous path from its questions to real-world utility (practical application, decision-relevant knowledge, or a concrete artifact), with the path traceable and the utility asserted in falsifiable form; (G2) declares where its premises end (derived claims vs. unanalyzable primitives or named imported inputs); and (G3) does not derive legitimacy solely from

membership in an academic discipline, abstraction level, or knowledge silo. Exploratory research remains valuable iff it maintains an open grounding path. The criterion is program-level and directional, not per-result.

1. Introduction

1.1 The problem

The research ecosystem allocates enormous resources — trillions in public and private funding, millions of careers, and now billions of compute-hours — to work whose connection to any human problem is untraceable. The problem is not that such work exists; exploration is necessary. The problem is that the allocation system rewards the *wrong signals*: disciplinary prestige cascades, citation counts that fail as proxies for quality outside narrow fields, and novelty claims unbounded by the depth of their premises.

The corpus in which this paper is written already diagnosed the failure modes. *Two-Faced Scientific Methodology* shows how theoretical attractor states resist falsification through systemic incentives — funding, peer review, publication — rather than through evidence (Research 2025b). *Institutional Reform* traces the \$35B quantum-computing bubble to a prestige cascade and VC timeline mismatch, proposing time-bound falsification registries and mandatory independent verification (Collective and Quni-Gudzinis 2026a). *The Meta-Pattern of Reification in Physics* documents the systematic mistaking of formalisms for reality — what it calls the “misallocation of intellectual resources” (Quni-Gudzinis 2026c). And the *So What of Knowledge* inquiry already posed the foundational question this paper answers: if the universe functions without human understanding, human knowledge must serve a purpose beyond comprehension — knowledge should enable action, not just description (Research 2025a).

What none of these supply is an **operational criterion** — a test that a researcher, a funder, or a self-auditing AI pipeline can apply to a research program and use to decide whether it deserves advancement, while remaining compatible with genuine exploratory research. That is the gap this paper fills.

1.2 The amplification problem

The urgency is compounded by automation. Science-of-science pipelines can now process grant-to-output linkages at unprecedented scale (Wang et al. 2025), and AI-driven systems are positioned to automate pattern discovery across the research ecosystem (Chen et al. 2025). AI adoption now spans virtually all 333 tracked research fields (Hajkowicz et al. 2023). The same capability that can *measure* traceable utility can also *generate* untraceable abstraction at arbitrary volume. A criterion without a falsifiable structure would be gameable by exactly the systems that most need it.

2. The Grounding Criterion

2.1 Definition

A research program *P* (a sustained line of inquiry with a defined question set, methods, and outputs) earns a claim to advancement if and only if:

- **G1 — Live grounding path.** *P* can exhibit a live, non-vacuous path from its questions to real-world utility: a practical application, decision-relevant knowledge, or a concrete artifact, such that (a) the path is traceable (each hop from question → method → output → use is specifiable), and (b) the utility is asserted in falsifiable form (the claim can be checked against a defined outcome).
- **G2 — Premise-depth disclosure.** *P* declares where its premises end: which claims are derived, which are unanalyzable primitives, and which are named imported inputs — in the form of an explicit premise chain (L0, L1, L2, ...) with the claimed depth stated.
- **G3 — No silo legitimacy.** *P* does not derive its claim to advancement solely from membership in an academic discipline, abstraction level, or knowledge silo. Disciplinary membership is *evidence of context*, never a substitute for G1 and G2.

2.2 Program-level, directional, non-vacuous

Three design choices require defense.

First, **program-level, not per-result**: the criterion is applied to a sustained program, not to each publication. A single exploratory paper need not itself exhibit utility; the program of which it is a part must maintain an open grounding path — a stated candidate route to utility with a plausible mechanism. This is what preserves genuine exploratory research while excluding indefinite utility-deferral.

Second, **directional**: the criterion asks *whether the program is moving toward a grounding path*, not whether it has arrived. This is consistent with the empirical finding that interdisciplinary researchers sacrifice short-run citation impact for long-run funding performance (Sun et al. 2021): directionality rewards exactly the persistence that silo-exit requires.

Third, **non-vacuous**: “all knowledge is eventually useful” is excluded by the falsifiable-assertion clause of G1. The program must name the utility it claims and the observation that would count against it. This addresses the tautology risk identified in the consilience literature (Quni-Gudzinis 2026b).

2.3 Why utility can be traced (and why current metrics fail)

The objection “utility is unmeasurable” is empirically false: grant-to-impact linkages are already operational at scale (Wang et al. 2025), and applied-research outputs are measurably “more easily distinguished and captured” than basic-research outputs (Holý and Šafr 2024). The deeper problem is that the metrics *in use* are the wrong proxies. Citation-based indicators systematically underestimate clinical and applied impact relative to basic research even within a single field (Eck et al. 2012); in the

arts, humanities, and most social sciences they are not effective quality proxies at all (Thelwall et al. 2022); and research-quality itself is fundamentally unobservable, so any indicator used causally risks invalidating itself (Traag 2022). The grounding criterion is designed to be *measured by traceable downstream impact, not by citation proxies* — precisely because the proxies are the silo’s self-report.

3. The Premise-Depth Disclosure Standard

A premise chain is a finite list $L_0, L_1, \dots, L_n$ such that:

- L_0 is an unanalyzable primitive (a value, an axiom, a starting point) — *asserted*;
- each L_i for $i \geq 1$ is either derived from earlier levels or is a *named imported input* (a result, framework, or assumption taken from outside the program);
- the claimed depth of the theory is n : the theory is exactly as deep as its premises, no deeper.

The standard requires three disclosures:

1. **The chain itself** — each level named, with derived-vs-imported status.
2. **The claimed depth** — an explicit “this theory goes as deep as L_n ” statement.
3. **The import list** — which premises are named imported inputs, and from where.

The standard extends the epistemic-scaffolding analysis of scientific progress (Quni-Gudzinis 2025) into a publication-time requirement, and it operationalizes the boundary-drawing discipline demonstrated in the configuration-space-topology program, where the spin-statistics connection was explicitly declared an imported input (from QFT) rather than a derived result (Quni-Gudzinis 2026a).

Canonical application — this paper: L_0 = value primitive (utility is a legitimate criterion); L_1 = empirical premise (silos cause low-traceable-utility research); L_2 = epistemic premise (premise disclosure increases testability). Claimed depth: as deep as L_0 – L_2 , nothing more. This paper proves no theorem; it proposes a criterion and a falsification design, and it is only as deep as its three premises.

4. The Anti-Pattern Taxonomy

Four anti-patterns recur across disciplines and are each detectable by the criterion:

Anti-pattern	Definition	Detection signal	Corpus evidence
Silo-lock	Legitimacy derives from disciplinary membership; cross-domain synthesis is penalized	G3 fails; no grounding path crosses the discipline boundary	Laws-of-Form 1969 → >50yr NEVER-connected (RES.009 silo table); ID work reviewed as “not real or too soft” (Hyrnsalmi et al. 2025); societal-AI research becoming <i>less</i> interdisciplinary (Markus et al. 2025)

Anti-pattern	Definition	Detection signal	Corpus evidence
Abstraction-for-its-own-sake	Formal elegance replaces utility; models reified as reality	G1 fails; no traceable use; reification pattern present	Reification meta-pattern (Quni-Gudzinis 2026c)
Utility-deferral-without-path	“This will be useful someday” with no mechanism, no timescale, no falsifiable assertion	G1 fails at the falsifiable-assertion clause	Gödel-incompleteness-as-excuse in <i>So What of Knowledge</i> (Research 2025a)
Premise-concealment	Derived claims presented as if derived from nothing; imported inputs unstated	G2 fails; claimed depth exceeds actual depth	Attractor-state resistance to falsification (Research 2025b)

The taxonomy is deliberately short. A longer taxonomy would become a checklist ritual — the very silo the criterion forbids (F3).

5. Falsification Design and the Traceable-Impact Audit

5.1 Falsification conditions

The criterion’s empirical content lives in its falsification conditions:

- **F1 (tests L1):** Conduct a traceable-impact audit over a defined population of research programs, stratified by grounding status (grounded vs. siloed-abstraction-only). If siloed programs produce real-world utility (defined as downstream application, decision adoption, or artifact use, measured via the Funding-the-Frontier-style linkage apparatus (Wang et al. 2025)) at rates statistically indistinguishable from grounded programs, L1 fails and the criterion loses its causal motivation.
- **F2 (tests L2):** Measure whether requiring premise-depth disclosure changes output quality or downstream utility. If disclosure measurably *reduces* either, L2 fails and the standard must be revised (e.g., disclosure-lite or deferred disclosure).
- **F3 (self-siloing meta-test):** Apply the criterion to the criterion. If its application reduces total utility — e.g., it becomes a compliance ritual that crowds out exploratory research, or it creates a new “grounding-gatekeeper” silo — the criterion is self-refuting and must be revised or withdrawn.

5.2 The audit procedure (G1 operationalized)

The audit has five steps: (1) enumerate the program's declared outputs; (2) link outputs to downstream entities via grant→paper→patent→policy→trial→news linkage data (Wang et al. 2025); (3) classify links by type (application, decision, artifact) and by traceability (complete path vs. dead-end); (4) score grounding status (grounded / partial / siloed) with the program's own premise chain (G2) as context; (5) compare utility rates across grounding strata for the F1 test, with field-normalization per the RAE-normalization literature (Kenna and Berche 2010) and the per-discipline evaluation methods (Rons et al. 2013).

The audit's value-detection layer can be partially automated: research-value classifiers over large abstract corpora are already operational (Jiang et al. 2025), and AI-driven science-of-science pipelines provide the pattern-discovery substrate (Chen et al. 2025). The audit remains human-auditable at the link-classification step — the criterion must not outsource its own falsification to an opaque model.

6. Relation to Prior Work

Within the corpus. This paper is the operational successor to the *So What of Knowledge* paradox (Research 2025a), the institutional diagnosis of *Institutional Reform* (Collective and Quni-Gudzinis 2026a), the falsification-resistance analysis of *Two-Faced Scientific Methodology* (Research 2025b), the reification critique (Quni-Gudzinis 2026c), and the epistemic-scaffolding quantification (Quni-Gudzinis 2025). It extends the consilience program (Quni-Gudzinis 2026b, 2026d, 2026a) from synthesis to *criterion*, and it treats the joules-per-solution metric (Collective and Quni-Gudzinis 2026b) as a canonical positive case: a research program whose grounding path is live (metric → procurement → benchmark) and whose utility is asserted in falsifiable form.

Externally. The criterion synthesizes the science-of-science measurement literature: funding concentration produces diminishing marginal returns (Mongeon et al. 2016), macro funding-output indicators are discipline-dependent (Leydesdorff and Wagner 2009), interdisciplinary collaboration yields higher gains in a substantial minority of field-pairs (Abramo et al. 2018), dedicated ID funding programs can be evaluated (Rons 2013a, 2013b), and interdisciplinary researchers attain better long-run funding performance (Sun et al. 2021). The two book-length anchors frame the compatibility claim: *Pasteur's Quadrant* established that use-inspired basic research is not a compromise but a category of its own (Stokes 1997), and *The New Production of Knowledge* documented the shift toward mode-2, context-driven knowledge production (Gibbons et al. 1994).

Compatibility, not opposition, to basic research. The criterion does not demand that every program be applied. It demands that every program be *able to state* its route to utility. Stokes' quadrant shows the strongest basic research is often use-inspired (Stokes 1997); the criterion makes that inspiration legible. Exploratory research with an open grounding path is fully advanced under G1–G3.

7. Implementation Procedure for Research Pipelines

The criterion is designed to be installed as a gate at research-project initialization (Phase 0), analogous to the standing governance gates already used in this pipeline (So-What, premise-depth, falsification — the paper’s own P0 required all three in writing before Phase 1 began).

Phase 0 grounding declaration (new-project gate): 1. **G1 declaration:** state the program’s candidate grounding path in one sentence a non-specialist can parse; name the utility claim and its falsifiable assertion. 2. **G2 declaration:** write the premise chain $L_0 \dots L_n$ with derived/imported status per level; state the claimed depth. 3. **G3 declaration:** name the disciplines/abstraction levels that contextualize the work; state explicitly that none of them *alone* justifies advancement. 4. **Exploration clause:** if the program is exploratory, state the open grounding path (candidate route to utility) that keeps it alive. 5. **Re-check gate:** re-run the declaration at each milestone; a program whose grounding path closes is flagged for review, not automatically terminated (F3 discipline).

The same procedure applies to AI-assisted pipelines, with an additional requirement: the pipeline must be able to *disclose its own premise depth* (what it imported, what it derived, what it asserts) — the epistemic-legibility requirement that prevents automation from becoming a silo-fodder amplifier.

8. Discussion and Limitations

Limitations (disclosed per UIA phases 3–4): (i) L1’s empirical base is currently case-based and corpus-anchored, not a controlled audit; the F1 audit design is the remedy, not the result. (ii) The criterion’s terms — “real-world utility,” “non-vacuous path,” “traceable” — require operational definitions that the audit procedure supplies but that remain contestable; the contest is the point. (iii) The paper is written from within a specific research culture (a solo-scientist, Zenodo-first, AI-assisted pipeline); the criterion must survive contact with institutional research cultures it does not share. (iv) Research-quality unobservability (Traag 2022) bounds any utility measurement; the criterion treats its own measurements as indicators, not truths.

What this paper does not claim: that utility-maximization is derivable from first principles (L0 is a value, not a theorem); that all abstraction is bad (only closed-path abstraction fails G1); that per-result utility is required (the criterion is program-level); that the criterion is easy to apply (it is a discipline, like premise-depth itself).

Disconfirmation conditions (standing invitation): any demonstration that a siloed, abstraction-only program of substantial scale produced traceable utility at grounded-program rates (F1), any demonstration that premise disclosure harms research output (F2), or any demonstration that the criterion itself has become a ritual that crowds out exploration (F3) counts as a disconfirmation event. The author invites adversarial validation of all three, and commits to revise or withdraw the criterion on the first replicated F1–F3 failure.

9. Conclusion

The research ecosystem does not need another philosophy of research; it needs a test. The grounding criterion — live path, declared premises, no silo legitimacy — is that test: short enough to state in one sentence, strong enough to fail (F1–F3), and honest about its own depth (as deep as L0–L2, no deeper). Its purpose is not to police exploration but to make the *grounding path* of every research program — including this one, including the pipeline that produced it — explicit, traceable, and falsifiable. Research with a purpose is not a constraint on science; it is the condition under which science can be held accountable for the enormous resources it claims.

Acknowledgments

The author acknowledges the QNFO research corpus and the standing editorial directives that motivated this work: the So-What gate (every artifact must answer “why should a reader care?”), the depth-of-premises test (a theory is only as deep as its premises), and the requirement that research maintain real-world utility and purpose while remaining compatible with exploratory inquiry.

References

- Abramo, Giovanni, Ciriaco Andrea D’Angelo, and Flavia Di Costa. 2018. *Do Interdisciplinary Research Teams Deliver Higher Gains to Science?* <https://arxiv.org/abs/1810.12637>.
- Chen, Renqi, Haoyang Su, Shixiang Tang, et al. 2025. *AI-Driven Automation Can Become the Foundation of Next-Era Science of Science Research.* <https://arxiv.org/abs/2505.12039>.
- Collective, QNFO Research, and Rowan Brad Quni-Gudzinis. 2026a. *Institutional Reform: Fixing the Epistemic Incentive Structure of Post-Classical Computing.* <https://doi.org/10.5281/zenodo.21299211>.
- Collective, QNFO Research, and Rowan Brad Quni-Gudzinis. 2026b. *The Joules-Per-Solution Metric: Definition, Measurement Protocol, and Anti-Gaming Provisions for Honest Computational Benchmarking.* <https://doi.org/10.5281/zenodo.21637028>.
- Eck, Nees Jan van, Ludo Waltman, Anthony F. J. van Raan, Robert J. M. Klautz, and Wilco C. Peul. 2012. *Citation Analysis May Severely Underestimate the Impact of Clinical Research as Compared to Basic Research.* <https://arxiv.org/abs/1210.0442>.
- Gibbons, Michael, Camille Limoges, Helga Nowotny, Simon Schwartzman, Peter Scott, and Martin Trow. 1994. *The New Production of Knowledge: The Dynamics of Science and Research in Contemporary Societies.* SAGE Publications. <https://doi.org/10.4135/9781446221853>.
- Hajkowicz, Stefan, Conrad Sanderson, Sarvnaz Karimi, Alexandra Bratanova, and Claire Naughtin. 2023. *Artificial Intelligence Adoption in the Physical Sciences, Natural Sciences, Life Sciences, Social Sciences and the Arts and Humanities: A Bibliometric Analysis of Research Publications from 1960-2021.* <https://arxiv.org/abs/2306.09145>.
- Holý, Vladimír, and Karel Šafr. 2024. *Are Economically Advanced Countries More Efficient in Basic and Applied Research?*

- <https://arxiv.org/abs/2405.05227>.
- Hyrnsalmi, Sonja M., Grisha Liebel, Ronnie de Souza Santos, and Sebastian Baltes. 2025. *Not Real or Too Soft? On the Challenges of Publishing Interdisciplinary Software Engineering Research*. <https://arxiv.org/abs/2501.06523>.
- Jiang, Hang, Tal August, Luca Soldaini, Kyle Lo, and Maria Antoniak. 2025. *Automatic Detection of Research Values from Scientific Abstracts Across Computer Science Subfields*. <https://arxiv.org/abs/2502.16390>.
- Kenna, Ralph, and Bertrand Berche. 2010. *Normalization of Peer-Evaluation Measures of Group Research Quality Across Academic Disciplines*. <https://arxiv.org/abs/1006.3863>.
- Leydesdorff, Loet, and Caroline Wagner. 2009. *Macro-Level Indicators of the Relations Between Research Funding and Research Output*. <https://arxiv.org/abs/0911.1044>.
- Markus, Dror Kris, Fabrizio Gilardi, and Daria Stetsenko. 2025. *Societal AI Research Has Become Less Interdisciplinary*. <https://arxiv.org/abs/2506.08738>.
- Mongeon, Philippe, Christine Brodeur, Catherine Beaudry, and Vincent Lariviere. 2016. *Concentration of Research Funding Leads to Decreasing Marginal Returns*. <https://arxiv.org/abs/1602.07396>.
- Quni-Gudzinis, Rowan Brad. 2025. *Quantifying Structural Unification and Epistemic Scaffolding in Scientific Progress*. <https://doi.org/10.5281/zenodo.17229528>.
- Quni-Gudzinis, Rowan Brad. 2026a. *A Consilient Gap Synthesis of the QNFO/QWAV Research Portfolio*. <https://doi.org/10.5281/zenodo.21782596>.
- Quni-Gudzinis, Rowan Brad. 2026b. *The Consilience Framework: From Valuation Theory to the Void — a Cross-Domain Synthesis*. <https://doi.org/10.5281/zenodo.21804073>.
- Quni-Gudzinis, Rowan Brad. 2026c. *The Meta-Pattern of Reification in Physics*. <https://doi.org/10.5281/zenodo.19605445>.
- Quni-Gudzinis, Rowan Brad. 2026d. *Ultrametric Consilience Atlas: Cross-Domain Applications of p-Adic Mathematical Structure*. <https://doi.org/10.5281/zenodo.21722395>.
- Research, QNFO. 2025a. *So What of Knowledge: Reclaiming Purpose*.
- Research, QNFO. 2025b. *Two-Faced Scientific Methodology: How Theoretical Attractor States Resist Falsification*.
- Rons, Nadine. 2013a. *Interdisciplinary Research Collaborations: Evaluation of a Funding Program*. <https://arxiv.org/abs/1307.6784>.
- Rons, Nadine. 2013b. *Output and Citation Impact of Interdisciplinary Networks: Experiences from a Dedicated Funding Program*. <https://arxiv.org/abs/1307.6778>.
- Rons, Nadine, Arlette De Bruyn, and Jan Cornelis. 2013. *Research Evaluation Per Discipline: A Peer-Review Method and Its Outcomes*. <https://arxiv.org/abs/1307.7033>.
- Stokes, Donald E. 1997. *Pasteur's Quadrant: Basic Science and Technological Innovation*. Brookings Institution Press.
- Sun, Ye, Giacomo Livan, Athen Ma, and Vito Latora. 2021. *Interdisciplinary Researchers Attain Better Performance in Funding*. <https://arxiv.org/abs/2104.13091>.
- Thelwall, Mike, Kayvan Kousha, Mahshid Abdoli, et al. 2022. *Is Research Funding Always Beneficial? A Cross-Disciplinary Analysis of UK Research 2014-20*. <https://arxiv.org/abs/2212.05418>.

- Traag, V. A. 2022. *Science of Science – Citation Models and Research Evaluation*. <https://arxiv.org/abs/2207.11116>.
- Wang, Yifang, Yifan Qian, Xiaoyu Qi, et al. 2025. *Funding the Frontier: Visualizing the Broad Impact of Science and Science Funding*. <https://arxiv.org/abs/2509.16323>.