

The Universal Ignorance Audit: A Fifteen-Question Method for Systematic Inquiry into the Structure of Not-Knowing

Rowan Brad Quni-Gudzinis

2026-08-29

Abstract

The Universal Ignorance Audit is a fifteen-question, five-phase method for systematically interrogating the structure of not-knowing in any domain. Where conventional epistemic tools treat ignorance as an absence to be filled, the audit treats it as an active, structured state with architecture: hidden assumptions (scaffolds), representational confusions (map–territory errors), felt anomalies (wobbles), protected zones (taboo and identity-threatening questions), and productive capacities (actionable and relational ignorance). This paper presents the method’s development history, its theoretical grounding in agnotology, specified ignorance, epistemic humility, and calibration research, and the complete fifteen-question instrument with administration protocol. It reports the meta-audit – the audit applied to itself – which revealed systematic biases (an analytic, individualistic, extractive slant) and generated corrective sibling questions. Three worked applications are summarized: an ultrametric Majorana qubit framework, the braid-group approach to topological quantum computation, and the electron as a case study in renormalized ontology. The method is offered as a portable, content-independent instrument for research self-audit, AI-assisted research governance, and epistemic hygiene. Version 0.4 (29 August 2026) incorporates a second post-publication adversarial review: the administration protocol gains a stakeholder-presence step, counterweight questions for the power and somatic probes, temporal-risk probes (asymmetric loss and cost of waiting), an anti-rumination check, and an explicit termination

condition; the review's revised fifteen-question scaffold-detection protocol is included as Appendix A. The canonical fifteen-question instrument is unchanged. Version 0.4.1 (29 August 2026) is a post-red-team correction: the Appendix A mapping table now accounts for all fifteen canonical questions (Question 7's disposition and Question 15's correspondence are stated explicitly), and Section 3.4's probe enumeration matches the table; the instrument and the administration protocol are unchanged from v0.4.

Keywords: ignorance; agnotology; epistemic humility; methodology; self-audit; falsifiability; AI-assisted research

1. Introduction

The question “What do we not know?” is usually treated as a prompt for enumeration: list the gaps, then fill them. This framing mistakes the edge of a map for the shape of the territory. Not-knowing is not a void between islands of knowledge; it is an active structure with load-bearing members, blind corners, defended zones, and generative capacities. To inquire into it systematically requires instruments suited to its structure, not tools designed for the acquisition of facts.

This paper formalizes one such instrument: the **Universal Ignorance Audit**, a fifteen-question, five-phase deep-inquiry method developed through iterative dialogue between a human researcher and an AI assistant on 9 August 2026. The method emerged from a cascade of six seed questions (“What don't we know? What can we know? What can we know with what we don't know? What can we do with what we don't know? What else can we know that we don't? How can we know what we don't know?”), was expanded through twelve extracted universal meta-questions, refined into a thirteen-question v2.0, meta-audited against itself, and stabilized as the fifteen-question instrument presented here.

The audit is deliberately content-independent. Its questions operate on the structure of any epistemic state – personal dilemma, scientific theory, business strategy, research program, or institutional belief – rather than on the content of that state. Its design goal is not to reduce ignorance but to make it legible: to convert diffuse unease into named structure, and named structure into actionable inquiry. The fifteen-question instrument has been stable since v0.1; the administration protocol was extended in v0.4 (Sections 3.3–3.4), and the version history is given in Section 10.

2. Related Work

The audit stands at the intersection of several established literatures.

Agnotology, the study of culturally and institutionally produced ignorance, establishes that not-knowing is frequently manufactured, maintained, and politically load-bearing rather than naturally occurring (Proctor and Schiebinger 2008). The audit's power-analysis question (Question 9) operationalizes this insight at the individual scale. **Merton's specified ignorance** – the deliberate, disciplined articulation of what a research program does not know and why that boundary matters – provides the audit's stance that ignorance can be a professional instrument rather than a failure (Merton 1987). **Firestein's argument that ignorance drives science** supplies the constructive framing: productive research lives at the frontier of articulate not-knowing (Firestein 2012).

The audit also engages the cognitive and epistemic-humility traditions.

Kahneman and Tversky's program on judgment under uncertainty documents the systematic overconfidence that motivates the audit's falsifiability and inversion questions (Tversky and Kahneman 1974). **Tetlock's expert political judgment research** shows that calibrated, self-critical forecasters outperform confident experts, supporting the audit's insistence on explicit disconfirmation conditions (Tetlock 2005). **Intellectual humility** scholarship argues that accurate self-assessment of cognitive limitation is an epistemic virtue with measurable correlates (Whitcomb et al. 2015).

The known/unknown taxonomy popularized by Rumsfeld – known knowns, known unknowns, unknown unknowns – is a useful first cut that the audit deliberately goes beyond (Rumsfeld 2002). The taxonomy classifies ignorance by the knower's awareness; the audit classifies it by its *structure*: scaffolding, representational confusion, protected zones, somatic signals, and relational demands. A scaffold is not an unknown-unknown; it is a load-bearing assumption so successful that it has become invisible.

3. The Method

3.1 Design Principles

Four design principles govern the audit:

1. **Content independence.** Questions must apply to any X without modification.

2. **Depth ordering.** Questions progress from surface structure to protected depth, so that earlier answers scaffold later ones.
3. **Self-applicability.** The instrument must be applicable to itself; its recursive meta-question is load-bearing, not ornamental.
4. **Action closure.** The audit must terminate in actionable ignorance – a concrete “what can I do with this uncertainty right now” – rather than in resignation or abstraction.

3.2 The Five Phases and Fifteen Questions

The audit is organized into five phases, each addressing a distinct layer of epistemic structure.

Phase 1 – Surface the Structure

1. **Scaffold detection.** *What are the hidden assumptions and scaffolds holding this situation or belief up?* Targets the infrastructure of thought: the load-bearing premises so successful they have become invisible.
2. **Map–territory hygiene.** *Where might the map be mistaken for the territory?* Targets representational confusion: the model, metric, or symbol mistaken for the thing itself.
3. **Wobble probe.** *What is the wobble? Where is the tension, anomaly, or thing that does not fit?* Targets pre-cognitive signals: the felt place where the model does not balance.

Phase 2 – Stress-Test the Frame

4. **Inversion.** *What if the opposite were true?* A clean cut that exposes hidden dependence on the current polarity.
5. **Falsifiability test.** *What would a world look like in which this was false?* A richer exercise than “how would I know I am wrong”: it demands a concrete alternative world, preventing cheap escape clauses.
6. **Invariant extraction.** *If I changed the base, origin, center, or frame, what would remain the same?* Identifies what might be real versus what is a projection of a given frame.

Phase 3 – Multiply Perspectives

7. **Radical perspectival shift.** *How would this look to a radically different observer?* Extends beyond human social perspectives to non-human, future, or theoretical observers.
8. **Externalized ignorance.** *Who knows something about X that I do not, and what would they say is my biggest blind spot?* Converts ignorance from private to relational.

Phase 4 – Uncover the Hidden Forces

9. **Power analysis.** *Who benefits from the current framing? Who or what is silenced, excluded, or harmed?* Shifts the audit from epistemology to political epistemology.
10. **Protected ignorance probe.** *What is the most dangerous question I could ask about this – the one that threatens my identity, safety, or certainties?* Surfaces the taboo; no audit is complete without it.
11. **Somatic and tacit dimension.** *What does this confusion or ignorance feel like? Where does it live in the body? What does the feeling itself know?* Opens the embodied channel that conceptual auditing misses.

Phase 5 – Act

12. **Willful ignorance.** *What do I already know but am pretending not to know?* Names self-deception as a form of maintained ignorance.
13. **Actionable ignorance.** *What can I do with this uncertainty right now, without needing to resolve it?* The audit's praxis requirement.
14. **Relational ignorance.** *What does the unknown want from me?* Shifts from instrumental to relational orientation toward not-knowing.
15. **Recursive meta-question.** *What question am I not asking that I should be asking, given all of the above?* The fractal operator; the question that eats itself and generates the next pass.

3.3 Administration Protocol

The audit is administered as a structured written or dialogic exercise. The following protocol stabilizes results across administrators:

1. **State the target explicitly.** X must be named before questioning begins; an implicit target produces an unfocused audit.
2. **Answer every question.** If a question seems inapplicable, the administrator must explain why and then stretch to find its relevance. Skipping is forbidden; stretching is mandatory.
3. **Write answers down.** Oral audits lose the texture that the somatic and power questions depend on.
4. **Do not resolve during Phase 1–4.** The audit is not a problem-solving session; premature resolution forecloses deeper unknowing.
5. **Allow silence after Question 14.** The relational question requires receptive rather than analytic attention.
6. **Run the meta-question.** The fifteenth question's answer becomes the seed of the next audit pass, unless the termination condition (step 11) has been met.

The following five steps were added in vo.4 (29 August 2026) in response to a post-publication adversarial review:

7. **Stakeholder presence.** Before answering the power-analysis question (Question 9), name who else is affected by the target and who is absent from the decision. The audit defaults to a single-agent frame; this step is its explicit break. The review's obligation counterweight ("what would I do if I were responsible for others, not only myself?") is answered here, not deferred.
8. **Counterweights.** The power and somatic questions each carry a paired counterweight question, answered immediately after them: for Question 9, "what would a credible outsider – someone who does not share my frame – say is happening?" (so that power analysis does not slide into self-doubt); for Question 11, "is this feeling information, noise, or avoidance?" (so that feeling is not treated as evidence).
9. **Temporal risk probes.** Before the meta-question, answer "what is the cost of being wrong in each direction?" and "what is the cost of waiting? Is there a decision point or deadline?" The audit must not treat time as free; some windows close, and some downsides are asymmetric.
10. **Anti-rumination check.** Answer "am I asking more questions to avoid deciding, acting, or risking?" If the answer is yes, the next step is an action, not a question.
11. **Termination condition.** The audit terminates when (a) an honest pass through Questions 1–14 changes nothing about what should be done next, or (b) the cost of waiting exceeds the expected value of further inquiry. The recursion of Question 15 remains available, but termination – not recursion – is the default.

3.4 The 29 August 2026 Revision

This version was prompted by a second post-publication adversarial review, conducted on 29 August 2026 through AI dialogue and preserved in the author's research notes. The review identified three weaknesses in the instrument:

1. **A single-agent frame.** The audit interrogates one knower; it had no question that breaks the frame toward who else is affected, who is absent, and what is owed to others.
2. **Time-blindness.** The instrument priced neither delay nor asymmetric loss; nothing prevented indefinite inquiry dressed as depth.
3. **Unbalanced probes.** The power question could slide into self-doubt, and the somatic question could slide into treating feeling as evidence, because

neither carried a counterweight.

The review also proposed a revised fifteen-question scaffold-detection protocol with a built-in termination condition. The revision responds to all three weaknesses: steps 7–11 of the administration protocol (Section 3.3) add the stakeholder break, the counterweights, the temporal risk probes, the anti-rumination check, and the termination condition; the revised protocol is reproduced in Appendix A. The canonical fifteen questions of Section 3.2 are unchanged, because the calibration register (Section 8) and the worked applications (Section 5) are keyed to their numbering, and because the revised protocol keeps nine of the canonical probes – scaffold detection (Q1), inversion and falsifiability (Q4–Q5), externalized ignorance (Q8), power analysis (Q9), protected ignorance (Q10), the somatic and tacit dimension (Q11), willful ignorance (Q12), and actionable ignorance (Q13) – while omitting the wobble (Q3), map–territory (Q2), invariant-extraction (Q6), radical-perspectival-shift (Q7), and relational-ignorance (Q14) questions; Question 15 survives as R15, now carrying the vo.4 termination condition.

The review’s closing stance – that inquiry should no longer treat uncertainty as inherently sacred and is useful only when it changes perception, action, or obligation – is not adopted verbatim, because Section 6.4’s claim that not-knowing is a capacity rather than a deficiency is retained. What vo.4 adopts is the stance’s operational content: the anti-rumination check (step 10), the termination condition (step 11), and the obligation counterweight (step 7) are its published forms. The two positions are complementary: the gift of not-knowing is real, and it is finite.

4. The Meta-Audit: The Audit Applied to Itself

The audit was applied to itself on the day of its development. The meta-audit’s findings are part of the method’s specification, because they bound its claims:

- **The audit assumes ignorance is primarily generative rather than harmful.** This is a scaffold; for some domains (safety-critical decisions, clinical contexts), the assumption is dangerous.
- **The audit can create the illusion of having touched ignorance.** Listing a blind spot is not seeing it; the map–territory error applies to the audit itself.
- **The audit has an analytic, masculine, extractive slant.** Despite the somatic question, its structure is linear and penetrating; it underweights the receptive, surrendering, not-doing dimension of unknowing. A Zen master’s

critique would be that the audit uses the unknown to sharpen the sword of knowing.

- **The audit benefits the articulate, time-rich individual.** It may silence those whose survival depends on certainty, or who know in non-verbal ways inaccessible to a question list.
- **Reproducibility is low by design.** The audit is a qualitative method; independent administration to the same target by different auditors produces different answers because the instrument is sensitive to the auditor's own scaffolds. This is a feature of the method, not a defect, but it means the audit's outputs are not mechanically reproducible in the quantitative sense.
- **The audit has a single-agent frame.** Every question addresses one knower's structure; who else is affected, who is absent from the decision, and what is owed to others were outside the instrument's frame. This limitation was named by the 29 August 2026 review and is remediated by the stakeholder-presence step and the obligation counterweight (Section 3.3, steps 7–8).
- **The audit is time-blind.** The original protocol priced neither the cost of waiting nor the asymmetry of losses, so nothing structurally prevented indefinite inquiry dressed as depth. The temporal risk probes and the termination condition (Section 3.3, steps 9 and 11) supply the missing pricing.
- **Two probes are unbalanced without their counterweights.** The power-analysis question (Q9) can slide into self-doubt, and the somatic question (Q11) can slide into treating feeling as evidence, because neither carried a paired counterweight. vo.4 binds each to one in the administration protocol (Section 3.3, step 8).

These limitations motivated the sibling questions generated by the meta-audit (relational ignorance, temporal patience, discernment between generative and harmful ignorance, the gift of not-knowing), several of which were incorporated into the final fifteen-question instrument; the three added above were remediated by the vo.4 protocol extensions (Section 3.3, steps 7–11).

5. Worked Applications

Three applications from the development day illustrate the method's operation. Each is summarized at the level of its findings, not its full protocol. The three were selected because they are distinct in domain (quantum information theory, algebraic topology, and foundations of physics) and in epistemic risk profile (a design claim, a feasibility question, and an ontological commitment), so they

exercise the audit across the three main failure modes it targets: scaffolds, map–territory errors, and wobbles. They are drawn from the author’s own research program because the audit was developed through dialogue with that program on the same day; no claim is made that they constitute a representative sample of all epistemic states.

5.1 Ultrametric Majorana Qubit Framework

Applying the audit to a qudit/p-adic ultrametric quantum-error-correction paradigm surfaced the scaffold that the Bruhat–Tits tree structure couples to the Standard Model dynamically – a premise with no worked mechanism. The inversion question (“what if the tree architecture still respects the known error-correction tradeoff, just in a different guise?”) produced the audit’s most valuable output: a null-experiment design – search for a massive particle whose internal state is eerily immune to local electromagnetic noise within a frequency window – that would test the framework without requiring the full theory.

5.2 Braid Groups and Topological Quantum Computation

The audit exposed the gap between the mathematical braid group and the physical anyon: universality theorems apply to the algebra, but the physical realization depends on unverified assumptions about quasiparticle statistics in specific materials. The most dangerous question (“is topological quantum computation forever out of reach?”) was identified and held rather than resolved.

5.3 The Electron as Renormalized Ontology

Applying the audit to the electron, electricity, and the quantum domain surfaced the scaffold of mathematical realism – the assumption that terms in successful equations correspond to entities in the territory. The map–territory question located renormalization as a bookkeeping entry mistaken for a description of nature: the bare mass does not exist in the territory. The wobble was identified as the 10^{120} vacuum-energy discrepancy, a fault line rather than a wiggle. The invariant question produced the finding that worldline topology – braiding and linking numbers – survives frame shifts, which is why topological quantum computation works.

6. Discussion

6.1 Universality and Its Limits

The audit's universality claim is structural, not metaphysical. Its questions target the machinery of any epistemic state – supports, boundaries, social embeddedness, felt texture, forbidden zones – and are therefore portable. The meta-audit establishes the limits of that universality: the instrument is not culture-free, not gender-neutral in its epistemology, and not appropriate as a substitute for domain expertise or for the receptive modes of knowing it cannot reach.

6.2 Relation to Falsifiability

The audit is compatible with, and extends, the falsifiability tradition. Question 5 (the false-world test) is a strengthened falsifiability condition: it demands not just a disconfirming observation but a complete alternative world, closing the escape hatch of “the theory still works in some limit.” The audit adds what falsificationism omits: the social, somatic, and political structure of why a belief is held despite its disconfirmability.

6.3 The Recursive Structure

The fifteenth question is a fractal operator. Applied to any answer, it generates a new question; applied to itself, it generates the next audit. This recursion is the audit's engine. Its discipline changed with v0.4: an audit that terminates because nothing changed has been run honestly enough to act. The termination condition (Section 3.3, step 11) is the brake on the recursion, and it resolves the original tension between the audit's two closing demands – that it open new passes, and that it end in action. The recursion remains available; termination is no longer treated as a failure of honesty but as the audit's normal exit.

6.4 The Gift of Not-Knowing

The audit's final stance is that not-knowing is not a deficiency but a capacity. The relational-ignorance question (14), which surfaces the aporia of the unknown, reframes the discomfort of ignorance as an offering: what is lost when understanding arrives is the openness, humility, and vision that ignorance made possible. For a synthesist – one whose value lies in connecting domains rather than mastering one – not-knowing is not the price of the work but its raw material.

7. Conclusion

The Universal Ignorance Audit is a fifteen-question, five-phase instrument for making the structure of not-knowing legible. It is content-independent, depth-ordered, self-applicable, and action-closing. Its meta-audit establishes both its power and its limits: it is an analytic instrument with an analytic instrument's blind spots, and it must be used in the knowledge that it can manufacture the illusion of contact with the unknown. Used honestly, it converts the wobble of confusion into a compass, and the terror of groundlessness into a research program.

The method is released as an open methodology for researchers, educators, and anyone operating at the frontier of what they do not know. Its most important instruction is its last – *what question am I not asking?* – with the caveat added in v0.4 that the recursion carries an off switch: when an honest pass changes nothing about what should be done next, or the cost of waiting exceeds the expected value of further inquiry, the audit terminates in action (Section 3.3, step 11). The v0.4 revision (stakeholder presence, counterweights, temporal risk probes, anti-rumination check, termination condition), the v0.4.1 correction, and the full version history are recorded in Section 10; the fifteen questions themselves are unchanged since v0.1.

8. Calibration Register and Frontier Questions

The following register pre-registers the method's own testable claims. They are recorded here so that the method's effectiveness can be assessed independently of its authors' enthusiasm.

- **[CHECK: 2027] A structured application of the audit to a scientific research program surfaces at least one load-bearing assumption not previously named by the program's practitioners.**
Strength: [STRONG] | Status: [PENDING]. This claim is falsified if structured audits routinely recapitulate only already-articulated assumptions.
- **[CHECK: 2027] The audit's power-analysis question changes the priority ordering of at least one research program's next actions.**
Strength: [MODERATE] | Status: [PENDING]. Falsified if power analysis routinely produces no actionable output.
- **[CHECK: 2028] The audit, applied to AI-assisted research outputs, catches at least one category of error (scaffold, map–**

- territory, or protected-ignorance) that standard verification misses.** Strength: [STRONG] | Status: [PENDING]. Falsified if audits of AI-generated research add no information beyond existing quality gates.
- **[CHECK: 2028] At least one application of the audit outside its development context (different researcher, different domain) reports it as net-positive for research self-governance.** Strength: [WEAK] | Status: [PENDING]. Falsified if external applications are net-negative or inert.
 - **[CHECK: 2028] The vo.4 protocol extensions – the counterweights, the temporal risk probes, and the termination condition – measurably reduce rumination and time-to-action without measurably reducing scaffold discoveries, relative to the vo.3 protocol.** Strength: [MODERATE] | Status: [PENDING]. Falsified if audits using the extended protocol terminate no earlier, or surface fewer scaffolds, than audits using the vo.3 protocol.
 - **[CHECK: 2028] At least one external administration uses the revised scaffold-detection protocol (Appendix A) and reports its termination condition as the decisive improvement over the canonical form.** Strength: [WEAK] | Status: [PENDING]. Falsified if external use of the revised protocol is reported as net-neutral or net-negative on termination.

9. Declarations

Funding: This research received no external funding.

Conflicts of interest: The author declares no conflicts of interest.

Data availability: The development dialogue that produced the audit is preserved in the author’s research notes (Obsidian vault, 9 August 2026); the 29 August 2026 revision dialogue is preserved in the author’s research notes (Obsidian vault, notes/v1/2026/08/29/Update Qs.md); the full fifteen-question instrument is reproduced in this paper.

Code availability: Not applicable.

Author contributions: Rowan Brad Quni-Gudzinis conceived the audit’s seed questions, directed the expansion and meta-audit, and authored this paper.

Use of artificial intelligence: This paper was authored by the named human author with AI assistance. The Universal Ignorance Audit was developed through iterative dialogue between the human author and an AI assistant (a large language model — the DeepChat agent running a DeepSeek large language

model, accessed through the DeepChat research environment); the AI contributed the initial extraction of universal meta-questions, the v2.0 expansion, the meta-audit analysis, and drafting support for the applications and related-work sections. All substantive claims, the final instrument, and the calibration register were reviewed and approved by the human author. The specific model versions and prompt history are preserved in the author's research notes (9 August 2026) and are available on request for verification. The v0.4 revision integrates recommendations from a post-publication adversarial review conducted through AI dialogue (29 August 2026); the review's full text is preserved in the author's research notes, and every recommendation was selected, edited, and approved by the human author before incorporation. No AI-generated citation, author, or numerical claim appears in this paper without human verification.

Ethics approval: Not applicable.

Consent for publication: Not applicable.

Reproducibility statement: The audit is a qualitative method; its reproducibility claim is the instrument itself plus the administration protocol in Section 3.3. Independent administration to the same target by different auditors will produce different answers; this is a feature of the method's sensitivity to the auditor's own scaffolds, and is documented as a limitation.

10. Version History

- **v0.1** (10 August 2026; DOI 10.5281/zenodo.21878943; CC BY-NC-SA 4.0). Initial release.
- **v0.2** (10 August 2026; DOI 10.5281/zenodo.21878976; CC BY-NC-SA 4.0). Publication-metadata normalization (frontmatter DOI and status corrected); the text is otherwise identical to v0.1.
- **v0.3** (12 August 2026; DOI 10.5281/zenodo.21901984; CC BY 4.0). Corrections identified by the first post-publication adversarial review: Whitcomb et al. 2015 bibliographic metadata corrected to PPR 94(3): 509–539; Merton 1987 page range corrected to 13: 1–29 per the publisher; reproducibility limitation added to the meta-audit; selection criteria added for the worked applications; the Question 14 label aligned with the relational-ignorance definition; the AI disclosure tightened. The v0.3 record was deposited under CC BY 4.0 while its document frontmatter retained “CC BY-NC-SA 4.0”; the v0.3 CC BY grant remains in force.
- **v0.4** (29 August 2026; DOI 10.5281/zenodo.22158133; CC BY 4.0). This version was the v0.4 initial release. Incorporates the second post-

publication adversarial review (29 August 2026): administration-protocol steps 7–11 (stakeholder presence, counterweights, temporal risk probes, anti-rumination check, termination condition); new Section 3.4 documenting the revision; the meta-audit extended with the single-agent-frame and time-blindness limitations; Section 6.3 amended to reconcile recursion with termination; Appendix A reproducing the review’s revised fifteen-question scaffold-detection protocol with a mapping table; two new calibration-register entries; the declarations updated with the revision provenance; the frontmatter license aligned to the deposited record license (CC BY 4.0). The canonical fifteen-question instrument of Section 3.2 is unchanged.

- **vo.4.1** (29 August 2026; DOI 10.5281/zenodo.22159625; CC BY 4.0). Post-red-team correction, applied after the post-publication adversarial analysis gate returned a CONDITIONAL verdict with zero HIGH findings: the Appendix A mapping table now accounts for all fifteen canonical questions (Question 7 – radical perspectival shift – is stated as having no counterpart in the revised protocol, and R15 is identified as the canonical Question 15 with the vo.4 termination condition attached); Section 3.4’s probe enumeration is aligned with the table (nine probes kept, five dropped); the review’s closing stance on uncertainty is acknowledged in Section 3.4; Section 7 notes the termination condition as the recursion’s off switch; the meta-audit records the unbalanced-probes limitation and its counterweight remediation. The fifteen-question instrument and the administration protocol are unchanged from vo.4.

References

Firestein, Stuart. 2012. *Ignorance: How It Drives Science*. New York: Oxford University Press.

Merton, Robert K. 1987. “Three Fragments from a Sociologist’s Notebooks: Establishing the Phenomenon, Specified Ignorance, and Strategic Research Materials.” *Annual Review of Sociology* 13: 1–29.

<https://doi.org/10.1146/annurev.so.13.080187.000245>

Proctor, Robert N., and Londa Schiebinger, eds. 2008. *Agnotology: The Making and Unmaking of Ignorance*. Stanford, CA: Stanford University Press.

Rumsfeld, Donald. 2002. “DoD News Briefing – Secretary Rumsfeld and Gen. Myers.” United States Department of Defense, February 12, 2002.

Tetlock, Philip E. 2005. *Expert Political Judgment: How Good Is It? How Can We Know?* Princeton, NJ: Princeton University Press.

Tversky, Amos, and Daniel Kahneman. 1974. "Judgment under Uncertainty: Heuristics and Biases." *Science* 185 (4157): 1124–1131.
<https://doi.org/10.1126/science.185.4157.1124>

Whitcomb, Dennis, Heather Battaly, Jason Baehr, and Daniel Howard-Snyder. 2015. "Intellectual Humility: Owning Our Limitations." *Philosophy and Phenomenological Research* 94 (3): 509–539.
<https://doi.org/10.1111/phpr.12228>

Appendix A: Revised Scaffold-Detection Protocol (29 August 2026)

The following fifteen-question protocol was proposed by the 29 August 2026 adversarial review as a revision of the canonical instrument. It keeps the canonical scaffold, power, inversion, and somatic probes; adds observable grounding, stakes and stakeholders, evidence symmetry, asymmetric loss, the cost of waiting, obligation, and a built-in termination condition; and drops the map–territory, wobble, invariant-extraction, radical-perspectival-shift, and relational-ignorance questions (Questions 2, 3, 6, 7, and 14). It is offered as an alternate form for time-pressured or multi-agent settings. The canonical fifteen questions of Section 3.2 remain the instrument of record; the two forms share the administration protocol of Section 3.3.

R1. What exactly is claimed or happening, in observable terms? R2. What do I want, what is at stake, and who else is affected? R3. What assumptions make this look the way it looks? Which would be most uncomfortable to discover were false? R4. What would a credible outsider say is happening? R5. Who benefits from this framing? Who is excluded, harmed, or silenced? R6. What is the strongest evidence for my current view? The strongest against it? R7. What would change my mind? What would I have to admit if the opposite were true? R8. What is the cost of being wrong in each direction? R9. What is the cost of waiting? Is there a decision point or deadline? R10. What is the smallest honest action I can take before full certainty? R11. What would I do if I were not afraid? What would I do if I were responsible for others? R12. What is my body signalling – and is it information, noise, or avoidance? R13. What am I pretending not to know? What am I protecting? R14. Am I asking more questions to avoid deciding, acting, or risking? R15. Is it time to stop asking and live the answer? If not, what one question is next?

Mapping to the canonical instrument: R1–R3 correspond to Question 1 (scaffold detection), with added observable grounding and stakes; R4 to Question 8 (externalized ignorance); R5 to Question 9 (power analysis); R6–R7 to Questions 4–5 (inversion and falsifiability), with explicit evidence symmetry; R8–R9 and R14 are new in v0.4 (Section 3.3, steps 9–11), and R15 is the canonical Question 15 (recursive meta-question) with the v0.4 termination condition (Section 3.3, step 11) attached; R10 to Question 13 (actionable ignorance); R11 to Question 10 (protected ignorance), with the obligation counterweight; R12 to Question 11 (somatic and tacit dimension), with its counterweight; R13 to Question 12 (willful ignorance). Questions 2 (map–territory hygiene), 3 (wobble probe), 6 (invariant extraction), 7 (radical perspectival shift), and 14 (relational ignorance) have no counterpart in the revised protocol.